

# DVI-to-RGB (Sink) 1.7 IP Core User Guide

Revised February 2, 2017; Author Elod Gyorgy

## 1 Introduction

This user guide describes the Digilent DVI-to-RGB Video Decoder Intellectual Property. This IP interfaces directly to raw transition-minimized differential signaling (TMDS) clock and data channel inputs as defined in DVI 1.0 specs for Sink devices. It decodes the video stream and outputs 24-bit RGB video data along with the pixel clock and synchronization signals recovered from the TMDS link.

## 2 Features

- Connects directly to top-level digital visual interface (DVI) port
- 24-bit video (clocked parallel video data with synchronization signals) output
- Display Data Channel interface with built-in EDID ROM
- Resolutions supported: 1920x1080/60Hz down to 800x600/60Hz (148.5 MHz – 40 MHz)
- Selectable preferred resolution in EDID
- Xilinx interfaces used: IIC, vid\_io
- Digilent interfaces used: TMDS

## 3 Performance

The IP constrains TMDS\_Clk to 165 MHz, the maximum frequency outlined in DVI 1.0 specifications. However, depending on the actual FPGA part or speed grade, the maximum supported frequency might be lower. If the top-level design fails timing on pulse-width checks inside the IP instance, TMDS\_Clk needs to be (re)constrained to the maximum frequency supported on the project part. Check the part datasheet for  $F_{MAX\_BUFIO}$ , which is the most likely reason for failed timing. TMDS\_Clk should be constrained for  $F_{MAX\_BUFIO}/5$ . Consequently, this is the maximum pixel clock frequency supported on that FPGA family and speed grade.

IP quick facts	
Supported device families	Zynq®-7000, 7 series
Supported user interfaces	Xilinx®: IIC, vid_io Digilent: TMDS
Provided with core	
Design files	VHDL
Simulation model	VHDL Behavioral
Constraints file	XDC
Software driver	N/A
Tested design flows	
Design entry	Vivado™ Design Suite 2016.4
Synthesis	Vivado Synthesis 2016.4

## 4 Overview

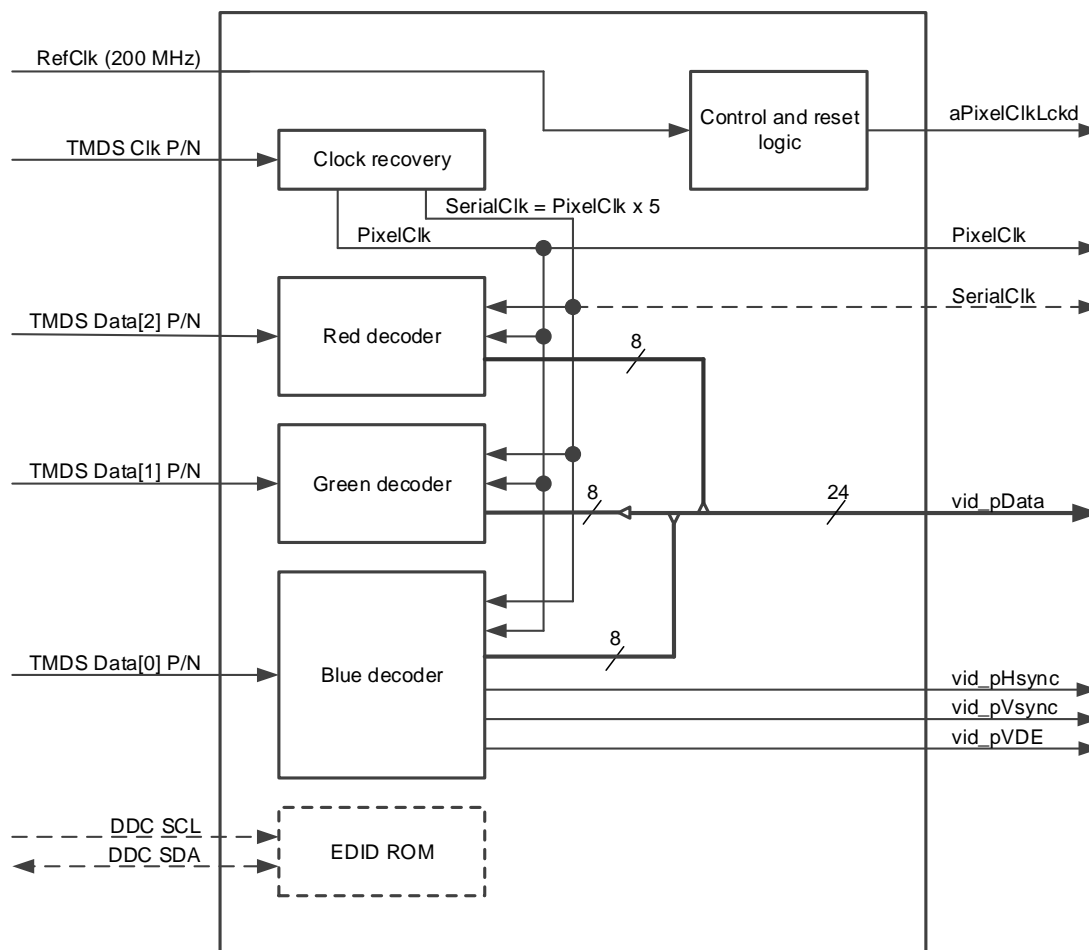


Figure 1. DVI to VGA converter block diagram.

The IP is built from multiple blocks: one clock recovery block, one data decoder block for each data channel (see [3], [4]), one optional DDC (Display Data Channel) block and one control/reset block.

### 4.1 Clock recovery

The clock channel carries a character-rate frequency reference. One character (or 10 bits) are transmitted every period on each data channel. Dedicated deserializer primitives, which require a fast serial clock, will be used to sample the serial data stream. The clock recovery block generates a serial clock and a pixel clock from the clock channel. The frequency ratio between the two clocks is 5:1.

Since the clock frequencies are relatively high and the recovered clocks have tight phase requirements, dedicated clocking primitives are instantiated inside this block. These can be seen in Figure 2. The MMCM primitive incorporates a voltage controlled oscillator (VCO) that has an operating range specified in the FPGA data sheet. Since there is no single set of MMCM parameters that maps the whole range of

DVI pixel clock frequencies to the VCO range, an IP customization parameter is available to optimize for the expected resolution and pixel clock frequency.

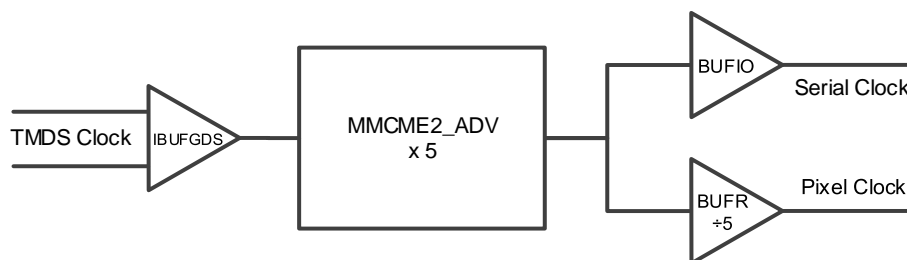


Figure 2. Clock network overview.

## 4.2 Data Decoder

There is no phase relationship between the recovered pixel clock or serial clock and the data channels. Furthermore, the data channels themselves are allowed a considerable skew between them. So the decoder block needs to align the phase of the serial clock with the data bit and find the character boundary in each data stream. Since this has to be done independently on each of the three data channels, it makes sense to vary the phase of the data streams while keeping the clock phase fixed.

### 4.2.1 Synchronization

To help with synchronization, the DVI protocol specifies period cues (control tokens) to be sent. These control tokens are sufficiently different from the rest of the data that their succession can be used for synchronization. Synchronization is automatically (re)started when a stable clock is detected and recovered. The time it takes for a lock to be achieved depends on the phase relationship of the clock and data streams. It should not last more than 1 minute.

DVI characters are 10-bits long, so a 10:1 deserialization of the data stream is needed. This can be achieved with two ISERDESE2 primitives in a cascaded DDR configuration. In this configuration, the master and slave ISERDESE2 take the serial data stream and sample it on both edges of a serial clock. Thus, for every five serial clock periods, ten data bits are sampled. This 10-bit data is then output synchronously with a divided clock, which is our pixel clock from the clock recovery block. Although this recovers 10-bit words from the data stream at a frequency which can be passed on to general logic inside the FPGA, it does guarantee the word actually starts at a character boundary or if the data stream is sampled when data is stable.

To find the best moment to sample the data stream (i.e., the middle of an open eye), an IDELAYE2 primitive is inserted in front of ISERDESE2. This primitive is capable of delaying the data signal in tap increments. In this IP, a 78ps increment is used for a total of 32 increments. For the highest pixel clock frequency supported (165 MHz), one bit period is covered in 7 tap increments. The goal is to find the tap delay value that shifts the data enough so that it gets sampled in the middle of its stable zone. The phase alignment module compares the 10-bit words with the four special control tokens. If a succession

of tokens is not recognized in a timeout period, we are in the jitter zone and it increments the tap delay. This is done repeatedly until control tokens are reliably recognized and the algorithm settles on the middle of the stable bit period (open eye). An “open eye” is defined by a succession of a minimum number of tap values (3) where the control token can be reliably detected, and it is delimited on both ends by a tap value where it cannot be. However, using this definition will miss open eyes that begin or end at the extremities of the tap delay range (0 or 31), because no two jitter zone delimiters will be found. So even if the open eye begins or ends in the extremities, if it is sufficiently long (16 tap increments), it will be considered a valid eye.

However, the IDELAYE2 primitive only provides a fine phase adjustment on the bit level, not covering the whole character. To find the character boundary, a coarse phase adjustment is needed. This is achieved by the “bitflip” feature of the ISERDESE2 primitive. If all the tap increments have been tried and control tokens are still not detected, it is assumed that we are not at the right character boundary. In this case, invoking “bitflip” causes either a shift right by one while bit or a shift left by three in the 10-bit word. After “bitflip” completes, phase alignment begins again, looping through the tap increments until tokens are found.

Phase alignment is considered completed when a succession of control tokens are reliably detected on all data channels. At this moment all three data channel are considered valid.

However, since inter-pair channel skew is not negligible and channels are aligned independently, the recovered data streams might have different delays on them. To eliminate this skew, the channels are bonded by buffering them in FIFO memories and holding them back independently until the video blanking period starts at the same time on all three. At this stage, all three data channels are valid and in-sync.

## 4.2.2 Decoding

The TMDS standard encodes data so that the serial data stream contains few transitions (0-to-1 or 1-to-0) and a DC balance (the same number of zeros and ones over a long time period). Every 10-bit character actually encapsulates 8 bits of useful data. The exception to this are the control tokens, which encapsulate 2 bits of control data. The data decoder block applies the decoding algorithm as specified in the DVI 1.0 specifications. After decoding we are left with control data in blanking periods or pixel data in active periods. Since each data channel carries one color, a 24-bit RGB pixel bus is output from the IP.

## 4.3 EDID ROM (Display Data Channel)

The DDC block emulates a read-only memory containing a default Diligent-branded extended display identification data (EDID). There are four variants available differing only in the preferred resolution. This parameter is user-selectable in the IP customization wizard. The EDID is defined in the src/\*\_edid.data files bundled with the IP (see section 6.4). Upon synthesis, the file corresponding to the preferred resolution gets incorporated into the netlist. Modifying these files is allowed, provided it remains compatible with DDC specs.

## 5 Port Descriptions

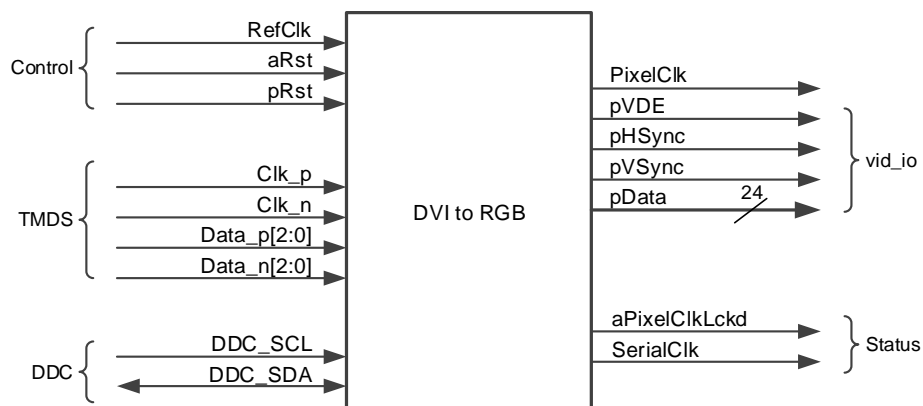


Figure 3. IP top-level diagram.

The signals of the DVI to VGA Core are listed and described in Table 1.

Signal Name	Interface	Signal Type	Init State	Description
RefClk	-	I	N/A	200 MHz reference clock.
aRst(_n)	-	I	N/A	Asynchronous reset of configurable polarity. Assert, if RefClk is not within spec.
pRst(_n)	-	I	N/A	Active-high reset synchronous with PixelClk. Configurable polarity.
Clk_p/Clk_n	tmcs	I	N/A	DVI Clock Channel.
Data_p[2:0]/Data_n[2:0]	tmcs	I	N/A	DVI Data Channel 0.
SCL	iic	I	N/A	Display Data Channel clock signal. Optional.
SDA	iic	IO	HiZ	Display Data Channel data signal. Optional.
PixelClk	rgb	O	N/A	Pixel clock recovered from the TMDS clock channel. Driven by BUFR.
pVDE	rgb	O	0	Video data valid: <ul style="list-style-type: none"> <li>• 1 = Active video.</li> <li>• 0 = Blanking period.</li> </ul>
pHSync	rgb	O	0	Horizontal synchronization video timing signal.
pVSync	rgb	O	0	Vertical synchronization video timing signal.
pData (23:0)	rgb	O	zeros	Video pixel data packed as RBG.

Signal Name	Interface	Signal Type	Init State	Description
aPixelClkLckd	-	O	0	Active-high asynchronous locked signal for PixelClk. When low, PixelClk is lost or not within specs.
SerialClk	-	O	N/A	Fast clock, toggling at five times the frequency of PixelClk. It is used internally for deserialization. Provided optionally as output for advanced use, like clocking an RGB2DVI core. Driven by BUFIO.

Table 1. Port descriptions.

## 6 Designing with the core

### 6.1 Constraints

The TMDS clock input Clk\_p/n is constrained in the IP to the maximum DVI clock frequency, 165 MHz. On some architectures this might result in timing impossible to meet. Depending on the application, if a lower pixel clock frequency is acceptable, the clock can be constrained on top-level, which will override the IP-internal constraints.

For example, to constrain the design for 720p resolution (74.25 MHz), calculate the clock period (13.468 ns), and add the following to a project XDC file to constrain the clock on the top-level input port:

```
create_clock -period 13.468 -waveform {0.000 5.000} [get_ports
hdmi_rx_clk_p]
```

### 6.2 Customization

The IP provides the following customizable parameters: the polarity of reset signals, PixelClk clock buffer type, the frequency range of TMDS clock, the preferred resolution to be declared in the bundled EDID, the availability of the DDC channel, and the serial clock output.

Enabling the DDC channel and serial clock will add the respective ports to the IP and are available to user logic.

The parallel pixel clock (PixelClk) is recovered by the use of a BUFR buffer. Since BUFR is restricted to a single clock region and the video data output from the core is synchronous to PixelClk, any downstream logic consuming video data is also restricted to this clock region. The option to re-buffer PixelClk introduces a BUFG after the BUFR and re-registers video data into the BUFG-domain. This will allow downstream logic to be placed anywhere on the device.

Setting the expected TMDS clock frequency enables the IP to instantiate FPGA primitives that respect timing requirements in the clock recovery logic. If the actual pixel clock recovered from the stream falls outside the range set here, the VCO operating range of the FPGA might not be respected and in extreme cases clock recovery might fail and the video stream will not be decoded properly.

The preferred resolution can be set, if the DDC channel is enabled. The resolution set here will select the proper initialization file for the emulated EDID ROM. This EDID will be read out by connected sources and *might* choose to transmit at this resolution.

### 6.3 Using SerialClk

While the fast serial clock is normally only used internally for deserialization, it is available and useful in other limited circumstances. The limitations arise from the fact that this clock has a frequency usually too large (five-times the PixelClk frequency) to clock user logic. For example, 1080p will result in a 742.5 MHz SerialClk. A BUFIO primitive drives this clock net to accommodate the large frequencies, which can only clock the I/O column in the same bank/clock region.

One case would be to use this clock to drive the Diligent RGB2DVI core, sharing the clocking logic between the two cores.

### 6.4 Bundled EDID

The IP comes with several EDID files each declaring a different preferred/native resolution. Files named \*\_edid.txt contain 128 bytes of EDID data in a format readable by Vivado synthesis. Each line has exactly one byte in binary format. Byte 0 is the first in the file. The following table summarizes the EDID in a human-readable format.

Monitor		Color characteristics	
Model name	Diligent DVI	Default color space	Non-sRGB
Manufacturer	DGL	Display gamma	2.20
Plug and Play ID	DGL0000	Red chromaticity	Rx 0.640 - Ry 0.330
Serial number	n/a	Green chromaticity	Gx 0.276 - Gy 0.594
Manufacture date	2014, ISO week 5	Blue chromaticity	Bx 0.145 - By 0.060
Filter driver	None	White point (default)	Wx 0.283 - Wy 0.297
EDID revision	1.4	Additional descriptors	None
Input signal type	Digital (DVI)		
Color bit depth	8 bits per primary color		
Color encoding formats	RGB 4:4:4		
Screen size	510 x 290 mm (23.1 in)		
Power management	Not supported		
Extension blocs	None		
DDC/CI	n/a		

Timing characteristics		Standard timings	640 x 480p at 60Hz - IBM VGA 800 x 600p at 60Hz - VESA 1024 x 768p at 60Hz - VESA 1680 x 1050p at 60Hz - VESA STD 1920 x 1080p at 60Hz - VESA STD 1280 x 1024p at 60Hz - VESA STD 1280 x 720p at 60Hz
GTF standard	Not supported		
Additional descriptors	None		
Preferred timing	No		
Detailed timing #1	1280x1024p at 60Hz (5:4)		

Table 2. Default EDID.

## 7 References

The following documents provide additional information on the subjects discussed:

1. Xilinx Inc., *UG471: 7 Series FPGAs SelectIO Resources*, v1.3, October 31, 2012.
2. Xilinx Inc., *UG472: 7 Series FPGAs Clocking Resources*, v1.6, October 2, 2012.
3. Xilinx Inc., *XAPP460: Video Connectivity Using TMDS I/O in Spartan-3A FPGAs*, V1.1, June 24, 2011.
4. Xilinx Inc., *XAPP495: Implementing a TMDS Video Interface in the Spartan-6 FPGA*, v1.0, December 13, 2010.
5. Xilinx Inc., *WP249: SPI-4.2 Dynamic Phase Alignment*, v1.3, July 6, 2011.
6. DDWG: *Digital Visual Interface DVI*, Revision 1.0, April 2, 1999.